

PhysioNet_Sepsis_Early_Detection

February 15, 2026

```
[52]: import pandas as pd

df = pd.read_csv("Dataset.csv", engine='python', on_bad_lines='skip')
print(df.head())
print(df.shape)
```

	Unnamed:	0	Hour	HR	O2Sat	Temp	SBP	MAP	DBP	Resp	EtCO2	...	\
0		0	0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	
1		1	1	65.0	100.0	NaN	NaN	72.0	NaN	16.5	NaN	...	
2		2	2	78.0	100.0	NaN	NaN	42.5	NaN	NaN	NaN	...	
3		3	3	73.0	100.0	NaN	NaN	NaN	NaN	17.0	NaN	...	
4		4	4	70.0	100.0	NaN	129.0	74.0	69.0	14.0	NaN	...	

	Fibrinogen	Platelets	Age	Gender	Unit1	Unit2	HospAdmTime	ICULOS	\
0	NaN	NaN	68.54	0	NaN	NaN	-0.02	1	
1	NaN	NaN	68.54	0	NaN	NaN	-0.02	2	
2	NaN	NaN	68.54	0	NaN	NaN	-0.02	3	
3	NaN	NaN	68.54	0	NaN	NaN	-0.02	4	
4	NaN	330.0	68.54	0	NaN	NaN	-0.02	5	

	SepsisLabel	Patient_ID
0	0	17072
1	0	17072
2	0	17072
3	0	17072
4	0	17072

[5 rows x 44 columns]
(1552210, 44)

```
[53]: import pandas as pd
import numpy as np

# Inspect basic info
df = df.copy()
print(df.shape)
print(df.head())
```

(1552210, 44)

	Unnamed: 0	Hour	HR	O2Sat	Temp	SBP	MAP	DBP	Resp	EtCO2	...	\
0	0	0	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	...	
1	1	1	65.0	100.0	NaN	NaN	72.0	NaN	16.5	NaN	...	
2	2	2	78.0	100.0	NaN	NaN	42.5	NaN	NaN	NaN	...	
3	3	3	73.0	100.0	NaN	NaN	NaN	NaN	17.0	NaN	...	
4	4	4	70.0	100.0	NaN	129.0	74.0	69.0	14.0	NaN	...	

	Fibrinogen	Platelets	Age	Gender	Unit1	Unit2	HospAdmTime	ICULOS	\
0	NaN	NaN	68.54	0	NaN	NaN	-0.02	1	
1	NaN	NaN	68.54	0	NaN	NaN	-0.02	2	
2	NaN	NaN	68.54	0	NaN	NaN	-0.02	3	
3	NaN	NaN	68.54	0	NaN	NaN	-0.02	4	
4	NaN	330.0	68.54	0	NaN	NaN	-0.02	5	

	SepsisLabel	Patient_ID
0	0	17072
1	0	17072
2	0	17072
3	0	17072
4	0	17072

[5 rows x 44 columns]

```
[54]: # Drop index-like and empty columns
drop_cols = ['Unnamed: 0', 'EtCO2', 'Unit1', 'Unit2']
df.drop(columns=drop_cols, inplace=True, errors='ignore')
print(df.columns)

Index(['Hour', 'HR', 'O2Sat', 'Temp', 'SBP', 'MAP', 'DBP', 'Resp',
       'BaseExcess', 'HCO3', 'FiO2', 'pH', 'PaCO2', 'SaO2', 'AST', 'BUN',
       'Alkalinephos', 'Calcium', 'Chloride', 'Creatinine', 'Bilirubin_direct',
       'Glucose', 'Lactate', 'Magnesium', 'Phosphate', 'Potassium',
       'Bilirubin_total', 'TroponinI', 'Hct', 'Hgb', 'PTT', 'WBC',
       'Fibrinogen', 'Platelets', 'Age', 'Gender', 'HospAdmTime', 'ICULOS',
       'SepsisLabel', 'Patient_ID'],
      dtype='object')
```

```
[55]: vital_cols = [
    'HR', 'O2Sat', 'Temp', 'SBP', 'MAP', 'DBP',
    'Resp', 'Platelets', 'Age'
]
label_col = 'SepsisLabel'
id_col = 'Patient_ID'

df = df[[id_col] + vital_cols + [label_col]]
print(df.head())
```

	Patient_ID	HR	O2Sat	Temp	SBP	MAP	DBP	Resp	Platelets	Age	\
0	17072	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	68.54	

```

1      17072  65.0  100.0   NaN    NaN  72.0   NaN  16.5      NaN  68.54
2      17072  78.0  100.0   NaN    NaN  42.5   NaN   NaN      NaN  68.54
3      17072  73.0  100.0   NaN    NaN   NaN   NaN  17.0      NaN  68.54
4      17072  70.0  100.0   NaN  129.0  74.0  69.0  14.0     330.0  68.54

SepsisLabel
0      0
1      0
2      0
3      0
4      0

```

```
[56]: # Forward-fill within each patient, then global median fill
df[vital_cols] = df[vital_cols].apply(pd.to_numeric, errors='coerce')
df[vital_cols] = (
    df.groupby(id_col)[vital_cols]
    .ffill()
    .bfill()
)
df[vital_cols] = df[vital_cols].fillna(df[vital_cols].median())
print(df.isna().sum())
```

```

Patient_ID      0
HR              0
O2Sat           0
Temp            0
SBP             0
MAP             0
DBP             0
Resp            0
Platelets       0
Age             0
SepsisLabel     0
dtype: int64

```

```
[57]: patient_labels = (
    df.groupby(id_col)[label_col]
    .max()
    .reset_index()
)
print(patient_labels[label_col].value_counts(normalize=True))
```

```

SepsisLabel
0      0.927311
1      0.072689
Name: proportion, dtype: float64

```

```
[58]: agg_features = {}
for col in vital_cols:
    agg_features[col + '_mean'] = (col, 'mean')
    agg_features[col + '_max'] = (col, 'max')
    agg_features[col + '_min'] = (col, 'min')
    agg_features[col + '_last'] = (col, lambda x: x.iloc[-1])

patient_features = (
    df.groupby(id_col)
    .agg(**agg_features)
    .reset_index()
)
print(patient_features.head())

   Patient_ID      HR_mean     HR_max     HR_min     HR_last  O2Sat_mean  O2Sat_max \
0            1  102.055556    117.0     76.0      84.0  91.398148    100.0
1            2   60.956522     94.0     54.0      55.0  97.086957    100.0
2            3   79.927083     93.0     68.0      78.0  95.333333     99.0
3            4  102.672414    113.0     93.0     108.0  98.155172    100.0
4            5   76.604167     88.0     61.0      80.0  97.677083     99.0

   O2Sat_min  O2Sat_last  Temp_mean ...  Resp_min  Resp_last  Platelets_mean \
0      85.0        85.0  36.636296 ...     17.0       18.0  322.833333
1      94.0        95.0  36.181739 ...      9.0       11.0  158.000000
2      91.0        97.0  37.469583 ...     17.0       26.0  476.354167
3      95.5        98.0  36.447931 ...     14.0       19.0  170.206897
4      96.0        97.0  37.072292 ...     14.0       16.0  276.562500

   Platelets_max  Platelets_min  Platelets_last  Age_mean  Age_max  Age_min \
0        338.0        317.0        338.0     83.14    83.14    83.14
1        158.0        158.0        158.0     75.91    75.91    75.91
2        488.0        465.0        486.0     45.82    45.82    45.82
3        220.0        144.0        144.0     65.71    65.71    65.71
4        288.0        273.0        288.0     28.09    28.09    28.09

   Age_last
0    83.14
1    75.91
2    45.82
3    65.71
4    28.09

[5 rows x 37 columns]
```

```
[59]: data = patient_features.merge(patient_labels, on=id_col)
print(data.shape)
print(data.head())
```

```
(40336, 38)

   Patient_ID      HR_mean    HR_max    HR_min    HR_last  O2Sat_mean  O2Sat_max \
0            1  102.055556  117.0     76.0     84.0  91.398148  100.0
1            2   60.956522  94.0     54.0     55.0  97.086957  100.0
2            3   79.927083  93.0     68.0     78.0  95.333333  99.0
3            4  102.672414  113.0     93.0    108.0  98.155172  100.0
4            5   76.604167  88.0     61.0     80.0  97.677083  99.0

   O2Sat_min  O2Sat_last  Temp_mean ...  Resp_last  Platelets_mean \
0       85.0        85.0  36.636296 ...     18.0    322.833333
1       94.0        95.0  36.181739 ...     11.0    158.000000
2       91.0        97.0  37.469583 ...     26.0    476.354167
3       95.5        98.0  36.447931 ...     19.0    170.206897
4       96.0        97.0  37.072292 ...     16.0    276.562500

   Platelets_max  Platelets_min  Platelets_last  Age_mean  Age_max  Age_min \
0         338.0        317.0        338.0     83.14    83.14    83.14
1         158.0        158.0        158.0     75.91    75.91    75.91
2         488.0        465.0        486.0     45.82    45.82    45.82
3         220.0        144.0        144.0     65.71    65.71    65.71
4         288.0        273.0        288.0     28.09    28.09    28.09

   Age_last  SepsisLabel
0     83.14          0
1     75.91          0
2     45.82          0
3     65.71          0
4     28.09          0
```

[5 rows x 38 columns]

```
[60]: from sklearn.model_selection import train_test_split
```

```
X = data.drop(columns=[id_col, label_col])
y = data[label_col]

X_train, X_test, y_train, y_test = train_test_split(
    X, y,
    test_size=0.2,
    random_state=42,
    stratify=y
)
print(y_train.value_counts(normalize=True))
```

```
SepsisLabel
0    0.927296
1    0.072704
Name: proportion, dtype: float64
```

```
[61]: from sklearn.linear_model import LogisticRegression

lr = LogisticRegression(
    max_iter=1000,
    class_weight='balanced',
    n_jobs=-1
)
lr.fit(X_train, y_train)
print(lr)

c:\Users\rohit\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\linear_model\_logistic.py:1184: FutureWarning: 'n_jobs' has no
effect since 1.8 and will be removed in 1.10. You provided 'n_jobs=-1', please
leave it unspecified.

warnings.warn(msg, category=FutureWarning)

LogisticRegression(class_weight='balanced', max_iter=1000, n_jobs=-1)

c:\Users\rohit\AppData\Local\Programs\Python\Python311\Lib\site-
packages\sklearn\linear_model\_logistic.py:406: ConvergenceWarning: lbfgs failed
to converge after 1000 iteration(s) (status=1):
STOP: TOTAL NO. OF ITERATIONS REACHED LIMIT

Increase the number of iterations to improve the convergence (max_iter=1000).
You might also want to scale the data as shown in:
https://scikit-learn.org/stable/modules/preprocessing.html
Please also refer to the documentation for alternative solver options:
https://scikit-learn.org/stable/modules/linear\_model.html#logistic-regression
    n_iter_i = _check_optimize_result()

[62]: from sklearn.metrics import (
    classification_report,
    roc_auc_score,
    confusion_matrix
)

y_pred = lr.predict(X_test)
y_prob = lr.predict_proba(X_test)[:, 1]

print(classification_report(y_test, y_pred))
print("ROC-AUC:", roc_auc_score(y_test, y_prob))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

          precision    recall  f1-score   support

          0       0.96      0.72      0.83     7482
          1       0.15      0.64      0.25      586
```

accuracy			0.72	8068
macro avg	0.56	0.68	0.54	8068
weighted avg	0.90	0.72	0.78	8068

ROC-AUC: 0.73285646644096

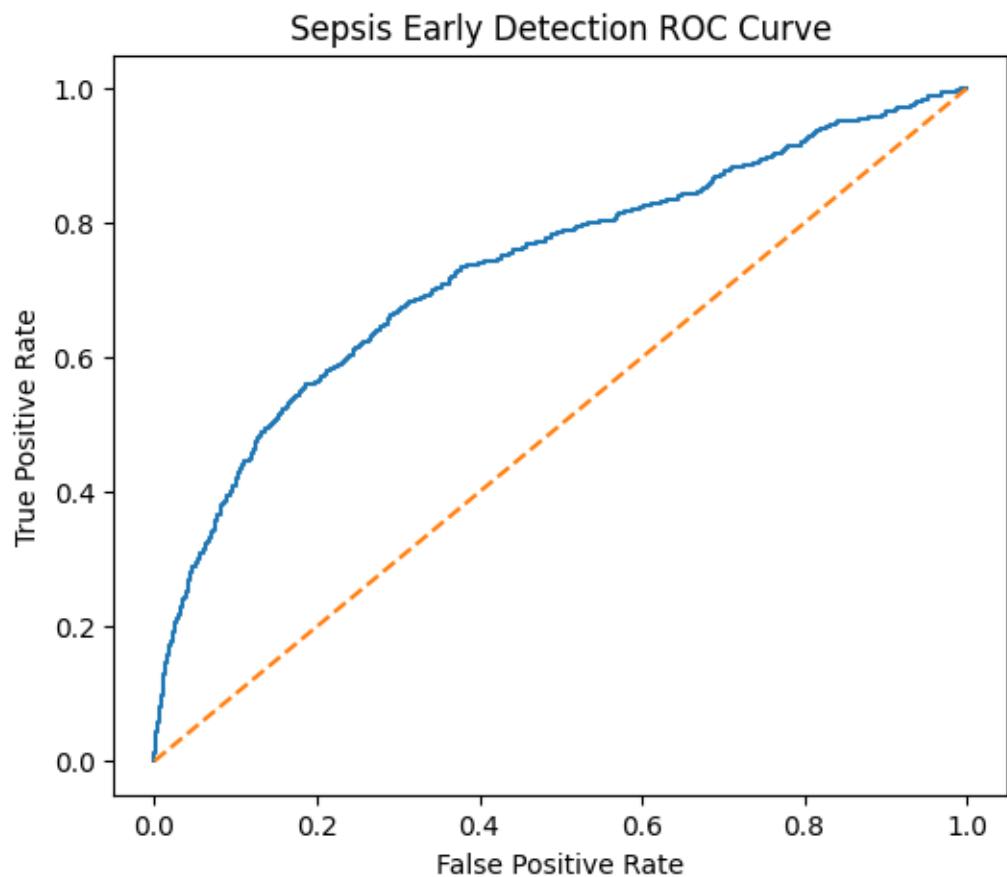
Confusion Matrix:

```
[[5406 2076]
 [ 210  376]]
```

```
[63]: import matplotlib.pyplot as plt
from sklearn.metrics import roc_curve

fpr, tpr, _ = roc_curve(y_test, y_prob)

plt.figure(figsize=(6, 5))
plt.plot(fpr, tpr)
plt.plot([0, 1], [0, 1], linestyle='--')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Sepsis Early Detection ROC Curve")
plt.show()
```



```
[64]: def risk_level(p):
    if p < 0.3:
        return "Low Risk"
    elif p < 0.6:
        return "Medium Risk"
    else:
        return "High Risk"
```

```
[65]: risk_df = pd.DataFrame({
    "Sepsis_Probability": y_prob
})
risk_df["Risk_Level"] = risk_df["Sepsis_Probability"].apply(risk_level)
print(risk_df["Risk_Level"].value_counts())
```

```
Risk_Level
Medium Risk      4292
Low Risk         2410
High Risk        1366
Name: count, dtype: int64
```

```
[66]: new_patient = pd.DataFrame({
    'Patient_ID': [999] * 6,
    'HR': [88, 90, 92, 95, 98, 102],
    'O2Sat': [97, 96, 95, 94, 93, 92],
    'Temp': [37.0, 37.2, 37.4, 37.6, 38.0, 38.4],
    'SBP': [120, 118, 115, 112, 108, 105],
    'MAP': [85, 82, 80, 78, 75, 72],
    'DBP': [75, 73, 70, 68, 65, 62],
    'Resp': [18, 19, 20, 22, 24, 26],
    'Platelets': [250, 245, 240, 230, 220, 210],
    'Age': [65] * 6
})
```

```
[67]: agg = {}
for col in vital_cols:
    agg[col + '_mean'] = (col, 'mean')
    agg[col + '_max'] = (col, 'max')
    agg[col + '_min'] = (col, 'min')
    agg[col + '_last'] = (col, lambda x: x.iloc[-1])

new_features = (
    new_patient
    .groupby('Patient_ID')
    .agg(**agg)
    .reset_index()
```

```
)  
X_new = new_features.drop(columns=['Patient_ID'])
```

```
[68]: new_prob = lr.predict_proba(X_new)[0][1]  
print("Sepsis Risk Probability:", round(new_prob, 3))  
print("Risk Level:", risk_level(new_prob))
```

Sepsis Risk Probability: 0.779
Risk Level: High Risk

```
[69]: from sklearn.ensemble import RandomForestClassifier  
  
rf_model = RandomForestClassifier(  
    class_weight='balanced',  
    random_state=42,  
    n_jobs=-1  
)  
rf_model.fit(X_train, y_train)  
print(rf_model)
```

RandomForestClassifier(class_weight='balanced', n_jobs=-1, random_state=42)

```
[70]: from sklearn.metrics import (  
    classification_report,  
    roc_auc_score,  
    confusion_matrix  
)  
  
y_pred_rf = rf_model.predict(X_test)  
y_prob_rf = rf_model.predict_proba(X_test)[:, 1]  
  
print("RandomForestClassifier Performance:")  
print(classification_report(y_test, y_pred_rf))  
print("ROC-AUC:", roc_auc_score(y_test, y_prob_rf))  
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred_rf))
```

RandomForestClassifier Performance:

	precision	recall	f1-score	support
0	0.93	1.00	0.96	7482
1	0.75	0.02	0.03	586
accuracy			0.93	8068
macro avg	0.84	0.51	0.50	8068
weighted avg	0.92	0.93	0.89	8068

ROC-AUC: 0.816240091121992

Confusion Matrix:

`[[7479 3]]`

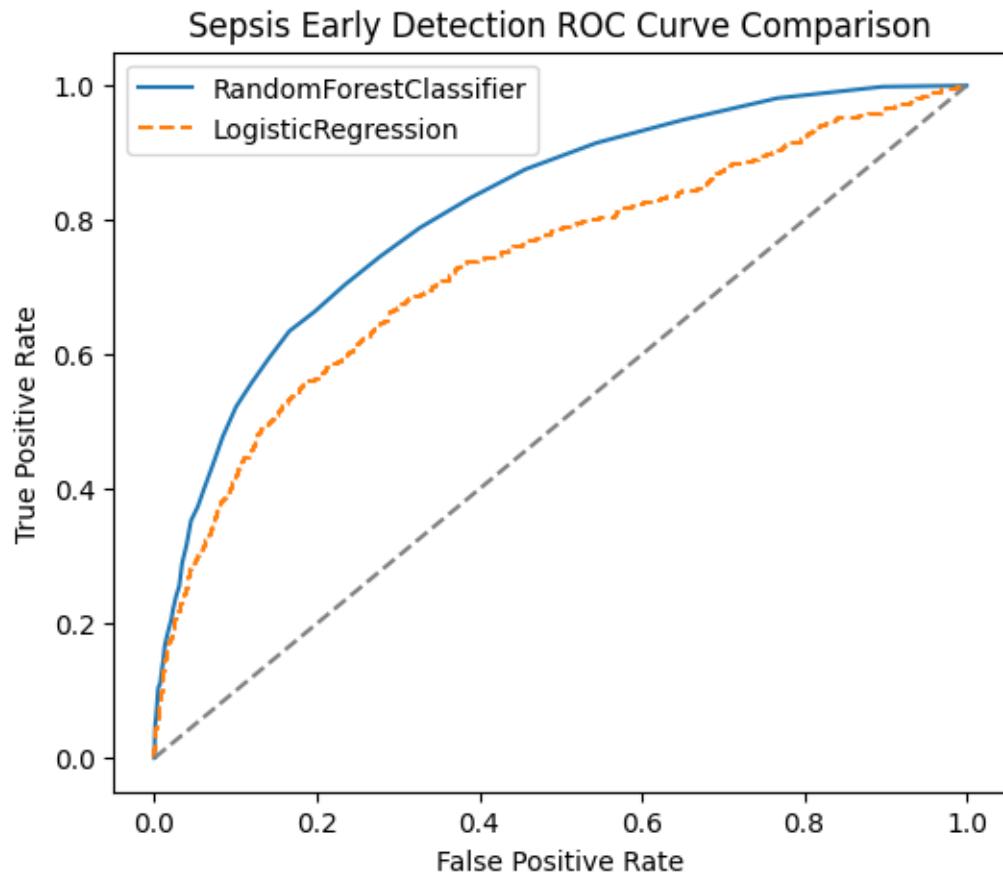
```
[ 577    9]]
```

```
[71]: from sklearn.metrics import roc_curve
import matplotlib.pyplot as plt

# Calculate ROC curve for RandomForestClassifier
fpr_rf, tpr_rf, _ = roc_curve(y_test, y_prob_rf)

# Calculate ROC curve for Logistic Regression
y_prob_lr = lr.predict_proba(X_test)[:, 1]
fpr_lr, tpr_lr, _ = roc_curve(y_test, y_prob_lr)

plt.figure(figsize=(6, 5))
plt.plot(fpr_rf, tpr_rf, label='RandomForestClassifier')
plt.plot(fpr_lr, tpr_lr, label='LogisticRegression', linestyle='--')
plt.plot([0, 1], [0, 1], linestyle='--', color='gray')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
plt.title("Sepsis Early Detection ROC Curve Comparison")
plt.legend()
plt.show()
```



```
[72]: new_prob_rf = rf_model.predict_proba(X_new) [0] [1]
      print("RandomForestClassifier Sepsis Risk Probability (New Patient):", ↪
            round(new_prob_rf, 3))
      print("RandomForestClassifier Risk Level (New Patient):", ↪
            risk_level(new_prob_rf))
```

RandomForestClassifier Sepsis Risk Probability (New Patient): 0.2
RandomForestClassifier Risk Level (New Patient): Low Risk